



Team Details

Team Name:

EdgeRestore-LK Stark Industries

SR. NO	ROLE	NAME	ACADEMIC YEAR
1	Team Leader	Vikaash R	3rd Year
2	Member 1	Mervin	2nd Year
3	Member 2	Nithilan T	3rd Year
4	Member 3	[ADD / REMOVE]	[ADD YEAR]

i A team can have up to 4 members including the team leader. Add rows if necessary.

COLLEGE NAME

Chennai Institute of Technology

TEAM LEADER CONTACT NUMBER

9600039825

TEAM LEADER EMAIL ADDRESS

r.vikaash301@gmail.com



Problem Statement Addressed



PS01 | Joint denoising + 2x grayscale super-resolution

WHY THIS MATTERS

Restore 128x128 noisy inputs to faithful 256x256 outputs. Preserve real edges and rare structures while suppressing speckle and Gaussian noise; hidden sources test generalization and speed.



Idea Description – EdgeRestore-LK



One deterministic model handles compound noise and 2x upscaling while prioritizing faithful structure over visually plausible hallucination.

KEY CONCEPT & APPROACH

6 residual groups x 6 residual channel-attention blocks (96 channels), zero-initialised head so the untrained network is exact bicubic.

SOLUTION OVERVIEW

PixelShuffle restores 2x detail; a bicubic skip anchors fidelity. Challenge-range outliers are retained.




Proposed Solution – Model and Training

LR $\rightarrow$ 6x[6 RCAB] channel-attention groups (96ch) $\rightarrow$ PixelShuffle $\rightarrow$ + bicubic.
Charbonnier + FFT-L1 training selects EMA iter 22000 by held-out validation PSNR.

VISUAL PROOF

Sample 000960 | PSNR 40.83 dB | SSIM 0.9744 | held-out filename-random split





Innovation and Empirical Advantage



A kernel-only controlled ablation improved PSNR, SSIM, and LPIPS; no GAN, diffusion, ensemble, or degradation router was added.

CONTROLLED 3x3 ->

5x5 ABLATION

PSNR 27.0756 -> 27.7637

SSIM 0.70638 -> 0.73490

LPIPS 0.36586 -> 0.32771

FINAL MODEL

6,939,289 parameters | 27.9 MB

29.35 ms model | 38.42 ms end-to-end

+6.111 dB PSNR over bicubic



Results




Quantitative metrics and visual evidence on the 320-image held-out split, identical protocol for every row.

QUANTITATIVE METRICS (320-image held-out)

PSNR: 29.0989 dB (bicubic: 22.9877 dB)
SSIM: 0.7784 (bicubic: 0.5193)
LPIPS: 0.2760 (bicubic: 0.4435, lower is better)
MS-SSIM: 0.9396 (bicubic: 0.7943)
+6.111 dB PSNR over bicubic baseline
0/320 images lose to bicubic on PSNR or MS-SSIM

VISUAL EVIDENCE: 000960 | degraded -> restored -> ground truth



PSNR 40.83 dB | SSIM 0.9744



Technology, Feasibility and Impact



A trained, deterministic PyTorch pipeline restores every released test array with a standalone command and frozen checkpoint.

IMPLEMENTATION AND TRUST BOUNDARY

All 400 official test outputs verified finite and within $[0,1]$ directly via the required inference.py entry point, 0 violations. Tests cover data, model, metrics, training and inference (257 tests). H100 and true OOD remain unverified.



PyTorch 2.11 | FP32



RTX 5070 Ti | 22,000 iters



+6.111 dB | 29.35 ms



GitHub & Video Link



Public Repository

Clone and run inference following the README - no manual edits required

GitHub: github.com/Vikaash301/kla-semicon

Contains: README.md, standalone inference.py (accepts positional or --input-dir/--output-dir args), training script, checkpoints/edgerestore_v2.pt, restored_test_outputs/ (400 images), requirements.txt, and requirements-freeze.txt (literal pip freeze).

Video: not provided (optional per template)



Research, Code and References



Research Basis + Reproducibility

Implemented evidence and submission links

2026 UCAN and ShiftLUT motivate wider spatial context; our 5x5 depthwise ablation is the measured adaptation.

GitHub: github.com/Vikaash301/kla-semicon Video: [ADD OPTIONAL DEMO URL]



References & Citations

Primary sources; full verified list in REFERENCES.md

UCAN – Expansive Receptive Fields for Lightweight SR (CVPR 2026)

ShiftLUT – Spatial Shift Enhanced LUTs for Restoration (CVPR 2026)

I4C Hackathon 2026 challenge and released paired dataset (i4c.in)